

▾ Multivariate Linear Regression

▾ Importing the libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets

# Load the diabetes dataset
diabetes = datasets.load_diabetes()

diabetes_df = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)
diabetes_df['target'] = diabetes.target

diabetes_df.head()
```

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	target
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0.019908	-0.017646	151.0
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0.068330	-0.092204	75.0
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194	-0.032356	-0.002592	0.002864	-0.025930	141.0
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0.022692	-0.009362	206.0
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0.031991	-0.046641	135.0

```
diabetes.feature_names

['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']
```

▾ Splitting the dataset into Features and Labels

```
# Split the dataset into features and labels
X = diabetes.data
y = diabetes.target
```

▾ Shuffling the dataset

```
def shuffle_data(X, y):
    # Convert input lists X and y into numpy arrays
    X, y = np.array(X), np.array(y)
    # Create an array of indices ranging from 0 to the number of rows in X
    indices = np.arange(X.shape[0])
    # Shuffle the indices randomly
    np.random.shuffle(indices)
    # Use the shuffled indices to shuffle the corresponding elements in X and y
    X, y = X[indices], y[indices]
    # Return the shuffled data as two arrays X and y
    return X, y
```

```
X, y = shuffle_data(X, y)
```

```
# make sure that the labels and features are still matching after shuffling the data.
print(X[:5])
print(y[:5])
```

```
[[ 0.01991321 -0.04464164 -0.02345095 -0.07108515  0.02044629 -0.01008203
  0.11859122 -0.0763945  -0.0425721  0.07348023]
 [-0.03094232 -0.04464164  0.04660684  0.01498661 -0.01670444 -0.04703355
  0.00077881 -0.00259226  0.06345592 -0.02593034]
 [-0.04183994 -0.04464164  0.12852056  0.0631868  -0.03321588 -0.03262872
  0.01182372 -0.03949338 -0.01599827 -0.05078298]
 [-0.08179786 -0.04464164 -0.0816528  -0.04009932  0.0025589  -0.01853704
  0.07072993 -0.03949338 -0.01090444 -0.09220405]
 [ 0.07076875 -0.04464164  0.06924089  0.03793909  0.02182224  0.00150446
 -0.03603757  0.039106  0.07763279  0.10661708]]
[ 92.  99. 259.  51. 220.]
```

▼ Splitting the dataset into the Training set ,Test set and dev set

```
from sklearn.model_selection import train_test_split

# Split the data into train and test sets, with test set size of 40%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

# Split the test set into dev and test sets, with test set size of 50% i.e 20% test, 20% dev
X_dev, X_test, y_dev, y_test = train_test_split(X_test, y_test, test_size=0.5, random_state=42)
```

```
# Print the shapes of the train, dev, and test sets
print("Train set: ", X_train.shape, y_train.shape)

print("Dev set: ", X_dev.shape, y_dev.shape)

print("Test set: ", X_test.shape, y_test.shape)
```

```
Train set:  (265, 10) (265,)
Dev set:   (88, 10) (88,)
Test set:  (89, 10) (89,)
```

```
# Number of features
n_features = X_train.shape[1]

# Initialize the weights and bias
weights = np.zeros(n_features)
bias = 0
```

```
def hypothesis(X, weights, bias):
    # Return the predicted value of y
    return np.dot(X, weights) + bias

def mean_squared_error(y_true, y_pred):
    # Return the mean squared error between y_true and y_pred
    return np.mean((y_true - y_pred)**2)

def cost(X, y, weights, bias):
    # Return the cost using mean squared error
    y_pred = hypothesis(X, weights, bias)
    return mean_squared_error(y, y_pred)
```

▼ Implementing the Gradient Descent Algorithm

```
def gradient_descent(X, y, weights, bias, learning_rate):
    """
    Perform a single iteration of gradient descent to update the weights and bias

    Parameters:
        X (np.array): input data with shape (n_samples, n_features)
        y (np.array): target data with shape (n_samples,)
        weights (np.array): weight coefficients with shape (n_features,)
        bias (float): bias term
        learning_rate (float): learning rate for gradient descent

    Returns:
        weights (np.array): updated weight coefficients with shape (n_features,)
        bias (float): updated bias term
    """
    # Calculate the predicted values using the current weights and bias
    y_pred = hypothesis(X, weights, bias)

    # Calculate the gradient for the weights
    d_weights = -2 * np.dot(X.T, (y - y_pred)) / X.shape[0]

    # Calculate the gradient for the bias
    d_bias = -2 * np.mean(y - y_pred)

    # Update the weights and bias using the calculated gradients
    weights -= learning_rate * d_weights
    bias -= learning_rate * d_bias

    return weights, bias
```

▼ Training the model

```
# defining the training function
def train_model(X, y, weights, bias, learning_rate, n_iters):
    costs = []

    for i in range(n_iters):
        # Perform a single iteration of gradient descent
        weights, bias = gradient_descent(X, y, weights, bias, learning_rate)

        # Calculate the cost over the entire training set
        cost_ = cost(X, y, weights, bias)

        # Print the cost every 100 iterations
        if i % 100 == 0:
            print("Cost at iteration {}: \t {}".format(i, cost_))

        # Append the cost to the costs list
        costs.append(cost_)

    return weights, bias, costs

# Testing different learning rates
def plot_learning_rates(X, y, weights, bias, n_iters):
    learning_rates = [0.001, 0.01, 0.1, 1]
    costs = []

    for learning_rate in learning_rates:
        # Train the model with the current learning rate
        weights, bias, cost_ = train_model(X, y, weights, bias, learning_rate, n_iters)

        # Append the cost of the current model to the list of costs
        costs.append(cost_)

    # Plot the costs for each learning rate
    plt.plot(costs[0], label="0.001")
    plt.plot(costs[1], label="0.01")
    plt.plot(costs[2], label="0.1")
    plt.plot(costs[3], label="1")
    plt.legend()
    plt.xlabel("Iterations")
    plt.ylabel("Cost")
    plt.show()

plot_learning_rates(X_train, y_train, weights, bias, 1000)
```



```

Cost at iteration 0:      29906.623169800954
Cost at iteration 100:   22041.156710373103
Cost at iteration 200:   16768.577368337723
Cost at iteration 300:   13233.321735934167
Cost at iteration 400:   10862.135422096346
Cost at iteration 500:   9270.919946769354
Cost at iteration 600:   8202.317233369438
Cost at iteration 700:   7483.889257729543
Cost at iteration 800:   7000.097098959382
Cost at iteration 900:   6673.52602314482
Cost at iteration 0:      6436.119990319524
Cost at iteration 100:   5952.564760531499
Cost at iteration 200:   5875.441347643607
Cost at iteration 300:   5807.60539151369
Cost at iteration 400:   5742.001081682859
Cost at iteration 500:   5678.433907382942
Cost at iteration 600:   5616.832456104412
Cost at iteration 700:   5557.129852462427
Cost at iteration 800:   5499.261590251746
Cost at iteration 900:   5443.165413922547

```

The optimal learning rate is 0.01 because it consistently provides stability as the number of iterations increases.

```

Cost at iteration 200:   4581.3645486409105

```

```

# Train the model for 1000 iterations
weights, bias, costs = train_model(X_train, y_train, weights, bias, learning_rate=0.01, n_iters=1000)

```

```

Cost at iteration 0:      3406.0116263125155
Cost at iteration 100:   3022.9587153227576
Cost at iteration 200:   3016.2166402566872
Cost at iteration 300:   3016.0910949597546
Cost at iteration 400:   3016.081910732348
Cost at iteration 500:   3016.074797609971
Cost at iteration 600:   3016.067745666411
Cost at iteration 700:   3016.060719156765
Cost at iteration 800:   3016.0537170588937
Cost at iteration 900:   3016.0467389724004
Cost at iteration 1000:  3017.0804300623467

```

```

# defining the prediction function
def predict_model(X, weights, bias):
    # Calculate the predicted values using the current weights and bias
    y_pred = hypothesis(X, weights, bias)

    return y_pred

```

```

15000

```

```

# Make predictions on the dev set
y_pred = predict_model(X_dev, weights, bias)

# Calculate the mean squared error
mse = mean_squared_error(y_dev, y_pred)

print("Mean Squared Error: {}".format(mse))

```

```

Mean Squared Error: 2323.31737269564

```

▼ Testing the model on the test set

```

# testing the model on the test set
y_pred = predict_model(X_test, weights, bias)

# Calculate the mean squared error
mse = mean_squared_error(y_test, y_pred)


print("Mean Squared Error: {}".format(mse))

```

```

Mean Squared Error: 3165.4205143477434

```



✓ 0s completed at 4:11 PM

